BAPIL Technical Blueprint, Version 1.0: Audit and Revision Report

Evidence check, issue register, budget reconciliation, quality scores and publication checklist

Prepared for Md Naim Hassan Saykat (Independent Researcher, Bangladesh)

22 October 2026

Contents

About This Report 1

1. Executive Review 2
- 1.1 Overall assessment . . . 2
- 1.2 Five strongest aspects . . . 2
- 1.3 Ten most important weaknesses (by severity) . . . 3
- 1.4 Principal obstacles to credibility . . . 3
- 1.5 Recommendation . . . 3

2. Issue Register 3

3. Claim Verification Matrix 5

4. Technical and Financial Validation 7
- 4.1 Architecture review . . . 7
- 4.2 Benchmark methodology and acceptance criteria . . . 7
- 4.3 Implementation maturity . . . 7
- 4.4 Budget reconciliation . . . 7
- 4.5 Financial assumptions and missing inputs . . . 8

5. Governance and Implementation Redesign 8

6. Change Log 8

7. Final Quality Assessment 9
- 7.1 Scores by category . . . 9
- 7.2 Separate dimensions . . . 10
- 7.3 Why not 10 . . . 10

8. Publication Readiness Checklist 10

Sources Consulted for This Audit 11

About This Report

This report records an audit of the pre-publication draft of the *Bangladesh AI Policy & Innovation Lab (BAPIL): Technical Blueprint, Version 1.0* (English and Bangla builds of 21 October 2026, 117 and 116 pages), and the revision that followed. The revised Blueprint (140 pages in English, 134 in Bangla) is Deliverable 6 of the review mandate and is not reproduced here; this report contains the other deliverables: the executive review (Section 1), the issue register (Section 2), the claim verification matrix (Section 3), the technical and financial validation (Section 4), the governance and implementation redesign (Section 5), the change log (Section 6), the final quality assessment (Section 7) and the publication readiness checklist (Section 8). The revision was made in two passes: the first applied the corrections in

the issue register; the second re-read every section, including those the first pass did not need to change, and fixed the remaining internal inconsistencies (issues I-35 to I-41). The Blueprint itself now carries its revision history and an unresolved-items register (Annex V, Tables 68 and 69).

Who carried out the audit. The audit was carried out with the assistance of an AI model (Claude), working to a multidisciplinary review mandate supplied by the author. It is not an independent external review by a panel of human experts, and it does not replace review by qualified Bangladesh legal counsel, a public-finance specialist or an independent evaluator. Where this report says a claim was "verified", it means the claim was checked against the source named, on the date given.

What was reviewed. The full English text; all 59 original tables and 15 figures; the budget model (`blueprint/src/budget_model.py`); the reference implementation, its tests and its evaluation output; the GitHub repository and its release page; the two existing Zenodo records; and the Bangla edition, compared with the English edition for numbers, dates, tables, headings and cross-references (by an automated section-by-section comparison and by reading the revised passages).

What was not reviewed. (1) Statutory texts. The official Laws of Bangladesh database (bdlaws.minlaw.gov.bd) refuses automated retrieval, so the gazetted texts and commencement sections of the 2025 ordinances and 2026 Acts could not be read; legal statements rely on press reports and secondary legal commentary, and are marked accordingly. (2) References 1 to 83, inherited from the White Paper, were not re-verified one by one; only those that support claims examined here were checked. (3) No full copy-edit of the Bangla prose was made; Bangla changes were limited to the revised passages, numerals and terminology consistency. (4) The reference implementation was not security-tested.

1. Executive Review

1.1 Overall assessment

The draft was a strong, unusually complete policy-design document: a clear problem, a credible institutional model, a careful public-data-first stance, an operational risk-tier and gate system, a working (if small) reference implementation, a long-term roadmap and a costed plan. Its weaknesses were not of concept but of evidence and status. It described several things as existing that did not yet exist (a public GitHub release and DOI-archived code), stated legal positions more firmly than its sources allowed (enactment treated as commencement; features of the 2025 ordinances attributed to the 2026 Acts; a labelling requirement attributed to the draft AI policy that the cited sources do not contain), described a proposed project as "under implementation", and presented a budget without escalation, separate contingency, recruitment ramp or scenarios. Its evaluation targets were sound but lacked sample sizes, decision rules and a precise definition of citation support. The plan lacked a list of the approvals and agreements that must precede month 1.

Score of the draft: 6.2 out of 10 (weighted; Section 7). **Score after revision: 7.8 out of 10.** A score of 9 or 10 is not justified while the statutory texts remain unreviewed, no institutional approval or funding exists, the benchmark has not been built and the budget inputs have not been validated. These are not writing problems; they need external action (Section 8).

1.2 Five strongest aspects

1. **A coherent, proportionate design.** Public data first, advice not decisions, no surveillance; small pilots before scale; a clear handover route to agencies (G4).
2. **Operational responsible-AI controls.** Risk tiers with explicit rules, implemented and tested in code; lifecycle gates with checklists; a transparency register; an ethics committee with a binding veto.
3. **A working reference implementation.** Dependency-free Python with 49 passing unit tests, cited extractive answers with abstention, a hash-chained audit log and reproducible evaluation; documents are rebuilt from source.
4. **Bilingual by design.** Bangla treated as a first-class language in requirements, evaluation, normalisation rules and publication, with a full Bangla edition.
5. **Honest about limits in places.** The draft already labelled demonstration results as non-generalisable, called thresholds proposals and warned that legal texts needed confirmation; the revision extends that discipline to the whole document.

1.3 Ten most important weaknesses (by severity)

#	Weakness	Severity
1	Reference implementation described as "published on GitHub and archived on Zenodo" in release v1.1.0, which did not exist publicly (the public repository shows releases v1.0 and v1.0.1 only)	Critical
2	Legal status overstated: enactment treated as commencement; Ordinance provisions (NRDEX, national source-code repository) attributed to the Acts; chief data officers attributed to the wrong instrument; the deferred commencement of some ordinance provisions not mentioned	Critical
3	AI Hub project described as "under implementation" with a "reported cost" of Tk 192.66 crore; the source describes a proposal awaiting approval	High
4	Budget: no escalation, no separate contingency, full staffing from month 1, no scenarios, no one-time and recurring split, no tax treatment; annual totals that appear not to add (rounding)	High
5	No list of preconditions and authorisations before month 1; January 2027 start presented without dependencies	High
6	Evaluation targets without sample sizes, confidence rules, out-of-scope test set or a precise definition of citation support	High
7	Draft national AI policy said to emphasise labelling of AI-generated content; the cited sources do not say so	Medium
8	Governance bodies described without legal status, appointment and removal rules, or relation to statutory bodies; ERRC members nominated by the ICT Division (independence risk)	Medium
9	Foundation-model path not explicit about the retrieval and reranking steps before adaptation, and no feasibility reasoning against training from scratch	Medium
10	Risk register without early-warning indicators, contingencies, residual ratings or review frequencies; no complaint and appeal route	Medium

1.4 Principal obstacles to credibility

- **Government:** no approvals, no funding and no host institution yet; legal obligations unconfirmed.
- **Academic:** evaluation targets not yet measured; reliance on press reports for legal facts.
- **Technical:** the reference implementation is a prototype; dense retrieval, the portal, authentication and OCR exist only as design.
- **Funding:** the budget is a model with plausible but unvalidated inputs (salaries, quotations, overhead, duty status).

1.5 Recommendation

Revise (done), then publish as a proposal; pilot only after approvals. Publication is appropriate once three conditions hold: the GitHub release v1.1.0 is public, the Zenodo DOI is minted and printed in the PDFs, and the documents carry the status statements added in this revision. Implementation should not be described as ready until the preconditions in Annex T are met.

2. Issue Register

Severity: Critical (misleads on project status or legal position), High (material to a decision), Medium (weakens credibility), Low (presentation). Status: Resolved (fixed in the revision), Partly resolved (fixed as far as evidence allows), Open (needs external action).

ID	Location (draft)	Issue	Category	Severity	Evidence	Correction	Status
I-01	Front matter "What already exists"; 1.4 item 10; 1.7; 11.1; 16.4; 21.4; Table 1	Code described as published on GitHub and archived with a DOI; release v1.1.0 not public	Implementation status	Critical	GitHub releases page and remote tags on 22 Oct 2026: v1.0 and v1.0.1 only; code exists on local branch `blueprint-v1.0`	Reworded as "released with this Blueprint"; status table added; Annex D rule that the DOI record is published only after the release exists	Partly resolved (needs push and release)
I-02	1.7; 21.4	"Core ideas already work"	Overclaim	High	Demonstration over 10 documents and 24 questions only	Reworded: mechanisms can be built and tested on a small scale; pilot tests usefulness	Resolved
I-03	11.7	CI "runs on every change"	Implementation status	Medium	Workflows exist locally; never run on GitHub	"Configured to run once the repository is updated"	Resolved

ID	Location (draft)	Issue	Category	Severity	Evidence	Correction	Status
I-04	11.x	No maturity vocabulary	Clarity	Medium	Prompt requirement; terms used loosely	Definitions of conceptual design, prototype, reference implementation, tested release, production added (Section 11.1)	Resolved
I-05	Repository	No security policy; dependency practice undocumented	Security	Low	No `SECURITY.md` in repository	`SECURITY.md` added; dependency rules stated in 11.7	Resolved
I-06	2.6; 14.1; Table 42	Enactment treated as commencement	Legal	Critical	PL&B (June 2026): Ordinance in force except ss. 23 and 31 to 46 [113]; 18-month window [114]	Enactment and commencement separated; "to be confirmed" and "stricter reading" rules	Partly resolved (needs counsel)
I-07	8.11; 10.8; Tables 42, 43	NRDEX and national source-code repository attributed to the 2026 Act	Legal attribution	High	Only the 2025 Ordinance description supports them [84]	Attributed to the Ordinance; "to be confirmed in the Act"	Resolved
I-08	8.8; Table 42	Chief data officers attributed to the data-management framework	Legal attribution	Medium	PL&B: CDO is s. 23 of the personal data protection ordinance, deferred [113]; TBS article does not mention CDOs	Re-attributed; deferral noted	Resolved
I-09	14.1	Ordinance named "National Data Governance Ordinance"	Accuracy	Low	Official database title: জাতীয় উপাত্ত ব্যবস্থাপনা অধ্যাদেশ, ২০২৫ (National Data Management); TBS uses "Governance"	Both names given	Resolved
I-10	14.1	February 2026 data-residency amendment ordinance omitted	Completeness	Medium	Digital Policy Alert: promulgated 5 Feb 2026 [115]	Added to 8.12, 14.1, Tables 42 and 43	Resolved
I-11	Table 42	PDPA passage cited to [86], which does not name it	Citation	Medium	Daily Sun lists NDM Bill and Cyber Security Bill, not PDPA	PDPA cited to [85, 116]	Resolved
I-12	13.8; Table 41; Annex N	Labelling of AI content attributed to the draft AI policy	Misattribution	Medium	[88], [89] mention measures against false or misleading AI content, not labelling	Labelling presented as BAPIL's own rule	Resolved
I-13	3.6; 17.2; Table 42; Summary; Memo	AI Hub "under implementation", "reported cost" Tk 192.66 crore	Government initiative status	High	TBS 13 Jun 2026: proposal awaiting approval [35]; Observer 26 Aug 2026: USD 13 m grant agreement signed [117]	Status corrected; comparison only for scale, with caveat	Resolved
I-14	Table 4; 2.2	2023 AI readiness rank cited to [26], the 2025 index	Citation	Low	Rank 82 (2023) is from [80]	Cited to [80]	Resolved
I-15	Table 4; 2.2	GCI "top tier" vague	Precision	Low	BGD e-GOV CIRT: Tier 1 role-model country [119]	Wording specific; official source added	Resolved
I-16	Chapter 17; Table 50	18.49 total vs annual figures adding to 18.50	Arithmetic	Medium	Unrounded draft total Tk 18.4868 crore; independent rounding	Explained; note on rounding; reconciliation Table 61	Resolved
I-17	Chapter 17	No escalation in Years 2 and 3	Financial	High	CPI 8.34% (Sep 2026) [118]	8% escalation on taka costs	Resolved
I-18	Chapter 17	Contingency merged with overhead (10%)	Financial	High	Overhead pays the host; contingency covers estimating risk; merging them hides the risk allowance	10% overhead plus 5% contingency	Resolved
I-19	Chapter 17	Full staffing assumed from month 1	Financial	Medium	Director appointed in months 1 to 3 (15.4)	75% average Year 1 occupancy	Resolved
I-20	Chapter 17	No scenarios, one-time and recurring split, or tax treatment	Financial	High	Not present	Table 62 (low, base, high; cost types); duty and VAT flagged as input	Resolved
I-21	Chapter 17	Budget inputs not validated	Financial	High	No quotations, host salary scale or overhead rate	Inputs required listed (Annex Q)	Open
I-22	Table 28; 9.4; 9.5	No sample sizes or decision rule; citation metric loosely defined; no out-of-scope set; recall could be read as answer quality	Evaluation	High	Not present	Annex R (Table 63); 9.4, 9.5 revised; Table 28 definitions	Resolved
I-23	9.4	Language parity gated on a set too small to resolve 5 points	Evaluation	Medium	About 150 per language in v0: interval about plus or minus 6 points	Parity reported in v0, gated from v1	Resolved

ID	Location (draft)	Issue	Category	Severity	Evidence	Correction	Status
I-24	9.7; Table 30	Foundation-model ladder implicit; no feasibility reasoning	Technical	Medium	Not present	Five-step ladder; compute and data estimates; minimum viable scope	Resolved
I-25	Chapter 5; Annex F; Table 57 D4	ERRC appointment by government nominees; no removal rule; legal status unclear	Governance	Medium	Independence risk	Independent nominations panel; removal for cause; status statements; Annex S	Resolved
I-26	Chapter 5; Annex E	Charter could be read as overriding law	Governance	Medium	No precedence clause	Article 13 and Section 14.7	Resolved
I-27	13.6	No complaint and appeal route; statutory remedies not mentioned	Accountability	Medium	Not present	Complaint, appeal and external routes added	Resolved
I-28	Chapter 15	No preconditions or authorisations before month 1	Implementation	High	Not present	Annex T (Tables 65 and 66); 15.1 revised	Resolved (content); Open (approvals)
I-29	Chapter 19	Risk register lacks indicators, contingencies, residual ratings and review frequency	Risk	Medium	Not present	Annex U (Table 67)	Resolved
I-30	Front matter	No statement of what exists versus what is proposed	Clarity	High	Not present	Status table; disclaimer extended	Resolved
I-31	Bangla edition	ASCII digits in lists, citations and quarter labels	Editorial	Low	Author's request	Bangla numerals via babel; citations converted; Gantt labels	Resolved
I-32	Zenodo record 21363967	Record licence shows MIT while the White Paper text states CC BY 4.0; affiliation shows Université Paris-Saclay	Metadata (outside the Blueprint)	Low	Zenodo record page, 22 Oct 2026	Recommend correcting the record metadata	Open
I-33	Statutory texts	Gazetted texts, commencement and rules not reviewed	Legal	Critical	bdlaws refuses automated retrieval	Legal review scope defined (14.3)	Open (needs counsel)
I-34	References 1 to 83	Not re-verified item by item	Citation	Medium	Inherited from the White Paper	Stated in Annex P	Open
I-35	Table 44 MS1; 15.4; Summary	MoU signed at month 3, although month 1 is now defined by the MoU	Internal consistency	Medium	Annex T definition of month 1	MS1 and first 90 days aligned; "before month 1" added to the Summary	Resolved
I-36	Annex E; 14.3	Charter adopted at MS1 "after legal review", but the review is due before MS2	Internal consistency	Medium	Table 44 sequence	Charter adopted at MS1 and confirmed after the legal review	Resolved
I-37	10.3; Table 31; Summary; Memo	One or two inference GPUs, while the budget buys two GPU servers in Year 1 and two more in Year 2	Internal consistency	Medium	Table 49 equipment	Capacity text and Table 31 matched to the budget	Resolved
I-38	Table 31	No adaptation compute in 2027 to 2028, though Stage 1 may start at month 13	Internal consistency	Low	Section 9.7	Allocation from month 13 if Stage 1 is approved	Resolved
I-39	Table 54	Time saved "user-reported", while Annex R measures it in crossover studies	Measurement	Medium	Annex R	KPI measured by time-to-evidence studies	Resolved
I-40	Table 18	Training targets (240 people in Year 1) not reconciled with the KPI of 100 officials and the budget	Internal consistency	Low	Tables 49, 54	Counting rule stated; specialist cohorts funded from the specialist line	Resolved
I-41	Tables 25, 48	OCR error-rate sampling unspecified; calendar dates assume a January 2027 start	Measurement; schedule	Low	Not stated	Notes added	Resolved

3. Claim Verification Matrix

Sources were consulted on 21 and 22 October 2026. "Corrected wording" is the substance now used in the Blueprint.

Claim (draft)	Location	Best available source and date	Status	Explanation	Corrected wording	Further verification
PDPA 2026 enacted April 2026	1.1, 2.6, 14.1, Tables 4, 42	Securiti, 29 Jun 2026 (Law 63 of 2026) [85]; Azam, SSRN, 14 Sep 2026 [116]; ConsentStack (enacted 10 Apr 2026; some provisions delayed 18 months)	Source-reported	No primary text seen; one aggregator gives 10 April and a deferred commencement	"Reported as enacted in April 2026 (Act 63 of 2026) … commencement to be confirmed"	Gazette and section 1 of the Act
NDMA 2026 passed April 2026	Same	Daily Sun, 10 Apr 2026 [86]; Azam [116]	Source-reported	Bill passed on the final day of the session, 10 Apr 2026	"Bill reported passed on 10 April 2026"	Gazette; commence-ment
Ordinance commencement	New	Greenleaf and Al Razi, PL&B, 1 Jun 2026 [113]; Khan (SCL), 23 Apr 2026 [114]	Source-reported	PDPO in force immediately except ss. 23 and 31 to 46; 18-month window from 6 Nov 2025 described	Added	Whether carried into the Act
NRDEX, national source-code repository	8.11, 10.8, Tables 42, 43	TBS, 10 Nov 2025 [84]	Verified for the 2025 Ordinance only	Act text not seen	"As described for the 2025 Ordinance; to be confirmed in the Act"	Act text
Data residency amendment	New	Digital Policy Alert [115]; TBS, 20 Apr 2026	Source-reported	Ordinance No. 23 of 2026, 5 Feb 2026	Added	Status after the Act
Cyber Security Bill passed 10 Apr 2026	14.1	Daily Sun [86]	Source-reported	Listed among 24 bills	Kept, qualified	Gazette
AI-content amendment passed 30 Jun 2026	14.1	Dhaka Tribune, 30 Jun 2026 [87]	Source-reported	Passed by voice vote; exact day not stated	"Reported passed by 30 June 2026"	Gazette
Draft AI policy: categories, oversight committee, fund Tk 200 to 250 crore	2.6, 17.2, Table 42	Daily Star, 26 Jan 2026 [88]	Source-reported	Confirmed; oversight committee to be created by an Act of Parliament	Kept; Act detail added	Final policy text
Draft AI policy emphasises labelling	13.8, Table 41, Annex N	[88], [89]	Not supported	Sources mention measures against false or misleading AI content	Labelling is BAPIL's rule	None
Drafting from Q4 2026 by joint team	1.1, 14.2	Daily Observer, 10 Aug 2026 [91]	Source-reported	Year inferred from article date	Kept	None
AI Hub Tk 192.66 crore, under implementation	3.6, 17.2, Table 42	TBS, 13 Jun 2026 [35]; Daily Observer, 26 Aug 2026 [117]	Partly incorrect	Proposal awaiting approval; grant agreement (USD 13 m, 2026 to 2029) later signed	Status corrected	Planning Commission approval
EGDI 2024 rank 100, 0.657	2.2, Table 4	UN DESA country data [120]	Verified	0.65696; 111 in 2022	Kept	None
GCI 2024 top tier	2.2, Table 4	BGD e-GOV CIRT [119]	Verified (official)	Tier 1 role model	Specific wording	None
GAIRI 2023 rank 82	2.2, Table 4	Sultana et al. 2026 [80]	Source-reported	Index report not checked	Cited to [80]	Index report
GPUs per department 2 to 55	Table 4	[80], Table 2	Verified against source	Estimates for five departments	Kept	None
Exchange rate Tk 123.65	Table 49	Bangladesh Bank, 19 Oct 2026 [110]	Verified	Weighted average rate	Kept	Update at costing study
Pay scale grade 9 Tk 44,000	17.5, Table 49	Financial Express, 31 Aug 2026 [109]	Source-reported	Cabinet approval; gazette not seen	Kept, status noted	Gazette
Inflation 8.34% (Sep 2026)	Table 49	Dhaka Tribune, 7 Oct 2026 [118]	Source-reported	BBS release reported	Added	BBS release
L40S about USD 7,500; 4-GPU server USD 47,000 to 49,000	Table 49	Modal, 12 Aug 2025 [111]	Source-reported (vendor)	Used as basis for USD 25,000 per 2-GPU server	Kept	Local quotations
Reference implementation published (v1.1.0)	Front, 1.4, 11, 16.4	GitHub, 22 Oct 2026	Future milestone	Not yet public	"Released with this Blueprint" plus Annex D rule	Confirm after push
Demonstration Recall@10 0.92; abstention 6 of 6; 49 tests	11.5	Build output, 22 Oct 2026	Research finding	10 documents, about 1,700 passages, 24 questions	Kept with limits	None

4. Technical and Financial Validation

4.1 Architecture review

The twelve-component architecture is sound and proportionate: a governed corpus, hybrid retrieval, an answer service that cites or abstains, an audit log and a governance workspace, with zero-trust zoning and hosting in Bangladesh for any non-public data. Three points were clarified in the revision rather than redesigned. First, the boundary between what is implemented (the prototype) and what is design (dense retrieval, portal, authentication, OCR, production security) is now stated in Section 11.1 and Table 37. Second, the evaluation rules now separate retrieval from answer quality, so the architecture's trust properties are measured rather than assumed. Third, the foundation-model path now begins with retrieval improvements before any adaptation. No change to the component model was needed.

4.2 Benchmark methodology and acceptance criteria

The revised protocol (Annex R, Table 63) fixes: BanglaPolicyQA v0 at 500 answerable questions split 40% public development and 60% private test, plus 200 out-of-scope questions; a decision rule (a threshold is met when the point estimate meets it and the 95% interval is no wider than plus or minus 5 points); citation precision judged at sentence level by two blinded annotators with reported agreement; and minimum samples for each measure. The arithmetic behind the sample sizes:

Measure	Expected value	Sample	Approximate 95% half-width
Recall@10, all test questions	0.85	300	4.0 points
Recall@10, one language	0.85	150	5.7 points
Citation precision	0.95	400 sentences	2.1 points
Correct abstention	0.90	150	4.8 points
P2 extraction accuracy	0.99	1,000 cells	0.6 points

Language parity therefore cannot be decided in v0 and becomes a release condition from v1 (about 600 test questions per language). The thresholds themselves (Recall@10 of at least 0.85, citation precision of at least 0.95) were not changed: they are reasonable for an advisory search tool, and nothing in the evidence justified moving them.

4.3 Implementation maturity

Element	Maturity
Bangla normalisation, BM25 search with bilingual expansion, cited extractive answers with abstention, risk tiers, register, evaluation metrics, audit log, CLI, JSON API	Reference implementation (research prototype), tested software release 0.1.0 with 49 unit tests; not deployed
Dense retrieval, reranking, generative rephrasing with verification	Conceptual design
Portal, authentication, OCR pipeline, production security controls	Conceptual design
BanglaPolicyQA	Protocol and guidelines only
Pilots P1 to P6	Proposed

4.4 Budget reconciliation

The draft total of Tk 18.49 crore and the displayed annual totals (4.96, 6.71 and 6.83, which add to 18.50) differed only through rounding: the unrounded draft total was Tk 18.4868 crore. The draft did, however, omit escalation and a separate contingency and assumed full staffing from month 1. The revised base case is computed by the same published model:

Step	36-month total (Tk crore)	Change
Draft budget before this review	18.49	
Year 2 and 3 amounts restated at Year 1 prices	18.29	-0.20
Explicit escalation of taka costs (8% a year)	19.79	+1.50
Year 1 recruitment ramp (75% of posts filled on average)	19.29	-0.50
Evaluation items added (out-of-scope set, citation judgements, annual system audit)	19.54	+0.25

Step	36-month total (Tk crore)	Change
Separate contingency of 5% of direct costs = revised base case	20.43	+0.89

The revised base case is Tk 4.73, 7.49 and 8.21 crore in Years 1 to 3 (Tk 20.43 crore, about USD 1.65 million). About 21% of direct costs are one-time. Scenarios: low Tk 17.93 crore, high Tk 25.25 crore. Single-factor sensitivity: staff costs plus or minus 20% change the total by about Tk 1.98 crore; full Year 1 staffing adds Tk 0.52 crore; equipment 25% dearer adds Tk 0.71 crore; 10% escalation instead of 8% adds Tk 0.41 crore; a 10% weaker taka adds Tk 0.28 crore.

4.5 Financial assumptions and missing inputs

All figures are calculations from stated planning assumptions, not quotations. Inputs still needed before any figure can be called a cost estimate in the procurement sense: the host's salary scale, increment rules and overhead rate; quotations for servers, GPUs and storage; the host's import duty and VAT status; data-centre charges; confirmation of in-kind national compute; agreed rates for annotators and trainers. No funding has been requested, approved or committed.

5. Governance and Implementation Redesign

The revision keeps the original institutional model (a university-anchored lab under an MoU, a Steering Board, an ERRC with a binding veto and a TAB) and makes it operationally specific:

- **Status and limits (Annex S, Table 64).** For each body: legal or administrative basis, membership and appointment, chair and secretariat, decision rights and their limits, reporting and transparency, and conflict-of-interest safeguards. None has statutory power; internal rules never override the law (charter Article 13; Section 14.7).
- **ERRC independence.** Members are selected through an open call by an independent nominations panel and confirmed by the Steering Board; government may propose candidates but not hold a majority; removal only for cause by two thirds of the Board after a hearing. Decision D4 was revised accordingly.
- **Relation to statutory bodies.** The ERRC is not the oversight committee proposed in the draft AI policy (to be created by an Act of Parliament); statutory authorities and courts keep their powers; overlaps with the host's governing body and ethics committee are resolved explicitly.
- **Accountability.** A complaint and appeal route with time limits, external remedies preserved; annual independent audit of systems.
- **Implementation readiness (Annex T, Tables 65 and 66).** Thirteen preconditions with the instrument or authority each needs and its current status; a definition of month 1; ten work packages with owners, preconditions, acceptance tests, budget lines, risks and gates; what needs legislation and what does not.
- **Risk management (Annex U, Table 67).** For all eighteen risks: an early-warning indicator, a contingency, a residual rating and a review frequency.
- **Monitoring and evaluation.** KPIs are labelled as proposed targets; quality indicators follow Annex R; independent evaluations at months 18 and 36 are retained.

6. Change Log

Location	Change	Reason	Source or calculation	Impact	Remaining uncertainty
Front matter	Status table; disclaimer and Overview revised; budget line added; document counts 21 annexes and 67 tables	Separate what exists from what is proposed	GitHub check; Annex P	Meaning: lowers claimed maturity	None
1.1, 1.4, 1.6, 1.7	Legal timing qualified; item 17 and method step 6 added; implementation wording	I-01, I-06	[84, 85, 86, 113]	Meaning	Commencement

Location	Change	Reason	Source or calculation	Impact	Remaining uncertainty
2.2, 2.6, Table 4	EGDI and GCI sources; GAIRI citation; legal summary rewritten	I-14, I-15, I-06	[80, 119, 120]	Accuracy	None
3.6	AI Hub status	I-13	[35, 117]	Accuracy	Approval status
5.3, 5.5	Legal status of bodies; precedence of law	I-25, I-26	Design	Governance	None
8.8, 8.11, 8.12, 10.8	Attribution of CDO, NRDEX, source-code repository; residency amendment; stricter-reading rule	I-07, I-08, I-10	[84, 113, 115]	Legal	Act text
9.4, 9.5, Table 28	Benchmark composition, splits, statistical rationale, decision rule; refined metric definitions; answer coverage and availability rows	I-22, I-23	Wilson intervals	Evaluation; small cost increase	None
9.7, Table 30	Five-step ladder; feasibility estimate; minimum viable scope	I-24	6ND approximation [121]	Technical	Governed Bangla text volume
11.1, 11.7; SECURITY.md	Maturity terms; CI wording; security and dependency practice	I-03, I-04, I-05	Repository	Status	None
12.3	Exit criterion tied to protocol; labelled as target	I-22	Annex R	Evaluation	None
13.6, 13.8	Complaints and appeals; labelling attribution	I-27, I-12	[89]	Accountability	None
14.1, 14.2, 14.3, new 14.7; Tables 41 to 43	Legal chronology, cautions, legal review scope, status of BAPIL's arrangements	I-06 to I-12	[84 to 89, 106, 113 to 116]	Legal	Needs counsel
15.1	Definition of month 1	I-28	Annex T	Schedule	Decision dates
16.4	Publication in future tense	I-01	GitHub	Status	None
17.1 to 17.5; Tables 49, 50	Budget revised (Tk 18.49 to 20.43 crore), features, scenarios, sensitivity, rounding note, scale comparison caveated	I-16 to I-21	budget_model.py; [118]	Cost +Tk 1.94 crore	Inputs unvalidated
18.3, 19.2	KPIs as targets; link to Annex U	I-29	Design	M&E	None
21.4; Table 57	Conclusion wording; D4 revised; note on authorisations	I-02, I-25	Design	Governance	None
Annexes D, E, F, K, L, N	Release rule; Article 13; ERRC appointment and status; out-of-scope set; glossary; labelling row	I-01, I-25, I-26, I-22, I-12	Design	Various	None
New Annexes P to U; Tables 60 to 67	Evidence status; budget reconciliation and scenarios; evaluation protocol; governance status; preconditions and work packages; risk monitoring	Prompt deliverables	As listed	Adds about 20 pages	As listed
References	[113] to [121] added	New evidence	As listed	Citation	None
Summary and Memo (EN, BN)	Status lines; legal qualification; budget figures; D4; targets wording	Consistency	Body	Consistency	None
Bangla edition	All revisions translated; Bangla numerals in lists, citations, quarter labels; acronyms KOICA and VAT	I-31; consistency	Automated EN and BN number comparison	Editorial	Full copy-edit not done
Second pass (Ch. 10, 15; Annex E; Tables 18, 25, 31, 44, 48, 54; Summary; Memo)	MS1 and first 90 days; charter sequence; GPU capacity; Stage 1 compute; KPI measurement; training counts; OCR sampling; calendar caveat	I-35 to I-41	Cross-checks between plan, budget and protocol	Consistency	None
New Annex V; Tables 68, 69; Document Information	Revision history and unresolved-items register inside the Blueprint; revision status row	Keep the change log and open items with the document	This report	Traceability	Items U1 to U14

Material deletions. None of substance. The unsupported phrases ("under implementation", "core ideas already work", "published on GitHub and archived with a DOI", the labelling attribution) were replaced, not silently removed.

7. Final Quality Assessment

7.1 Scores by category

Weights reflect what decision-makers and funders most need from a blueprint: accuracy, evaluation and finance carry the most weight. The overall score is the weighted mean.

#	Category	Weight	Draft	Revised	Evidence for the revised score	Remaining limitations
1	Strategic clarity and public value	8	8.0	8.5	Clear problem, model and theory of change; status statements	Value not yet demonstrated with users

#	Category	Weight	Draft	Revised	Evidence for the revised score	Remaining limitations
2	Factual and legal accuracy	12	5.0	7.0	Misattributions corrected; enactment and commencement separated; Annex P	Statutory texts unreviewed; secondary sources
3	Technical architecture	8	7.5	8.0	Boundaries of implemented versus designed stated; ladder for models	Most components are design only
4	Evaluation and reproducibility	10	6.0	8.5	Protocol with samples, rules, blinded judging; reproducible build	No benchmark exists yet
5	Responsible AI, privacy and security	8	7.5	8.0	Tiers, gates, complaints, stricter-reading rule, security policy	Controls untested in operation
6	Governance and accountability	8	6.5	8.0	Annex S; independent ERRC selection; precedence of law	Bodies not constituted
7	Financial transparency and feasibility	10	5.5	8.0	Model, reconciliation, scenarios, sensitivity, cost types	Inputs unvalidated; no funding
8	Implementation readiness	8	4.5	6.0	Preconditions, month 1 definition, work packages	No approvals, host or funding
9	Research quality and citation integrity	8	6.0	7.5	Claim matrix; citations corrected; primary sources where possible	Refs 1 to 83 not re-verified; press sources for law
10	Internal consistency	8	6.5	8.5	Automated EN and BN comparison; totals and cross-references checked	Minor wording differences between editions
11	Bengali language and editorial quality	6	7.5	8.5	Bangla numerals; terminology aligned; new content translated	No full native copy-edit
12	Publication readiness	6	5.0	7.0	Builds clean; status explicit; Annex D rule	Release, DOI and counsel pending
	Weighted overall	**100**	**6.2**	**7.8**		

7.2 Separate dimensions

Dimension	Assessment after revision
Conceptual quality	High (8.5): coherent, proportionate, well argued
Research and factual accuracy	Moderate to good (7): corrected, but legal facts rest on secondary sources
Technical design quality	Good (8): sound architecture and evaluation design; small prototype
Institutional feasibility	Uncertain (6): plausible model, no host, sponsor or approvals yet
Financial feasibility	Plausible (7): modest, transparent, scenario-tested; inputs unvalidated
Operational readiness	Low (4): nothing can start before the Annex T preconditions
Publication quality	Good (8): consistent bilingual build in LaTeX, PDF, PNG and Markdown

7.3 Why not 10

A 10 would require a document that is evidence-backed, independently verifiable and ready for its purpose. Four things prevent that today, and none can be fixed by editing: (1) the statutory texts and commencement provisions have not been reviewed by qualified counsel; (2) the reference implementation, release and DOI are not yet public; (3) no host, sponsor, approval or funding exists, so implementation readiness is necessarily low; (4) the budget inputs have not been validated against quotations and the host's rules. With (1), (2) and (4) done, a score of about 9 would be defensible for the document as a publication-ready proposal; (3) concerns the project, not the document.

8. Publication Readiness Checklist

Resolved in this revision

- Status of every artefact stated; maturity terms defined; claims classified (Annex P).
- Legal misattributions corrected; enactment and commencement separated; stricter-reading rule.
- AI Hub status corrected; labelling attribution corrected; citations corrected.
- Budget reconciled, revised and scenario-tested; rounding explained.
- Evaluation protocol with samples and decision rules.
- Governance status, appointment and limits; complaints and appeals; precedence of law.
- Preconditions, work packages and risk monitoring.
- Bilingual consistency checked; Bangla numerals; documents rebuilt; tests pass (49).

Partly resolved

- Reference implementation status: wording now accurate at publication, provided the release is published first.
- Bangla editorial quality: revised passages and numerals done; full copy-edit not done.

Requires external confirmation

- Push the branch and publish GitHub release v1.1.0; confirm it appears on the releases page.
- Reserve the Zenodo DOI, rebuild the PDFs with it, then publish the record (after the release).
- Correct the White Paper Zenodo record's licence field (MIT shown; text says CC BY 4.0) if that was unintended.

Requires legal counsel

- Read the gazetted PDPA 2026, NDMA 2026, Cyber Security Act 2026 and amendment, and the 2026 amendment ordinance; confirm enactment, publication and commencement dates, provisions in force, rules issued, and BAPIL's role (controller, processor or fiduciary). Then correct Tables 42 and 43.

Requires technical testing

- Build BanglaPolicyQA v0 and run the Annex R protocol; security-test the prototype before any use beyond demonstration.

Requires institutional approval

- Decisions D1 to D6; host selection and governing-body resolution; MoU; charter adoption; ERRC selection.

Requires funding or procurement confirmation

- Year 1 funding commitment; host salary scale and overhead rate; equipment quotations and duty status; national compute confirmation.

What remains before the Blueprint can credibly be called publication-ready: the GitHub release and the Zenodo DOI. **Before it can be called implementation-ready:** legal review, approvals, a host and funding, as listed in Annex T.

Sources Consulted for This Audit

The full references are in the Blueprint. Sources consulted for verification on 21 and 22 October 2026 include:

- GitHub repository and releases page: https://github.com/md-naim-hassan-saykat/bapil/releases
- Zenodo records: https://zenodo.org/records/21363967, https://zenodo.org/records/21441172
- Securiti, PDPA 2026 overview: https://securiti.ai/bangladesh-personal-data-protection-act-overview/
- Greenleaf and Al Razi, Privacy Laws & Business, June 2026: https://www.privacylaws.com/reports-gateway/articles/int201/int201bangladesh/
- Khan, Sovereign Chambers of Law, April 2026: https://bd-scl.com/insights/personal-data-protection-ordinance-2025-compliance.html
- Azam, SSRN, September 2026: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=7436740
- Digital Policy Alert, amendment ordinance 2026: https://digitalpolicyalert.org/change/18757-personal-data-protection-amendment-ordinance-2026-ordinance-no-23-of-2026
- The Business Standard, ordinances gazetted (Nov 2025): https://www.tbsnews.net/node/1281356
- Daily Sun, bills passed (10 Apr 2026): https://www.daily-sun.com/bangladesh/868072/91-bills-passed-in-parliament-to-convert-ordinances-into-laws
- Dhaka Tribune, cyber security amendment (30 Jun 2026): https://www.dhakatribune.com/bangladesh/parliament/414016/cyber-security-amendment-bill-2026-passes-in
- The Daily Star, draft AI policy (26 Jan 2026): https://www.thedailystar.net/business/economy/news/draft-ai-policy-targets-innovation-better-public-services-4090431
- The Financial Express, draft AI policy (29 Jan 2026): https://thefinancialexpress.com.bd/national/draft-ai-policy-tightens-state-oversight

- The Daily Observer, AI policy drafting (10 Aug 2026): https://observerbd.com/news/587314
- The Business Standard, AI hub proposal (13 Jun 2026): https://www.tbsnews.net/tech/govt-establish-ai-hub-tk19266cr-project-backed-koica-1462206
- The Daily Observer, KOICA grant agreement (26 Aug 2026): https://observerbd.com/news/590061
- UN DESA E-Government Knowledgebase, Bangladesh: https://publicadministration.un.org/egovkb/en-us/Data/Country-Information/id/14-Bangladesh
- BGD e-GOV CIRT: https://cirt.gov.bd/who-we-are
- Bangladesh Bank exchange rate: https://www.bb.org.bd/en/index.php/econdata/exchangerate
- Dhaka Tribune, inflation September 2026: https://www.dhakatribune.com/business/421534/inflation-rises-in-september
- The Financial Express, 2026 pay scale: https://thefinancialexpress.com.bd/national/cabinet-approves-2026-pay-scale-sets-basic-salaries-from-tk-20000-to-tk-156000
- Modal, L40S prices: https://modal.com/blog/nvidia-l40s-price-article