Table 28. Evaluation metrics and proposed pilot exit thresholds

Dimension	Metric	Definition	Proposed threshold
Retrieval	Recall@10	Share of questions whose gold passage appears in the top 10 results	at least 0.85
Retrieval	MRR@10	Mean reciprocal rank of the first relevant passage	at least 0.70
Retrieval	nDCG@10	Rank-weighted relevance against graded judgements	at least 0.75
Grounding	Citation precision	Share of cited passages that support the claim they are attached to	at least 0.95
Grounding	Unsupported-claim rate	Share of answer sentences with no supporting citation	at most 2%
Abstention	Correct abstention	Share of unanswerable questions on which the system declines to answer	at least 0.90
Language parity	Bangla and English gap	Difference in Recall@10 between Bangla and English question sets	at most 5 points
Robustness	Variant spelling recall	Recall@10 on romanised and variant-spelling Bangla queries	at least 0.75
Safety	Prompt-injection success	Share of red-team injection attempts that change system behaviour	at most 5%
Human evaluation	Usefulness and fluency	Mean rating by pilot users and Bangla editors on a 5-point scale	at least 4.0
Efficiency	Latency p95	95th-percentile response time for search and for answers	3 s and 10 s

Thresholds are proposals for Gate G2 and G3 decisions; the ERRC and TAB may tighten them for specific uses.

Source: BAPIL Technical Blueprint, Version 1.0 (2026). CC BY 4.0.