Table 28. Evaluation metrics and proposed pilot exit thresholds

Dimension	Metric	Definition	Proposed threshold
Retrieval	Recall@10	Share of answerable test questions with at least one gold passage in the top 10 results	at least 0.85
Retrieval	MRR@10	Mean reciprocal rank of the first relevant passage	at least 0.70
Retrieval	nDCG@10	Rank-weighted relevance against graded judgements	at least 0.75
Grounding	Citation precision	Share of answer sentences whose cited passage fully supports the sentence, judged by two trained annotators	at least 0.95
Grounding	Unsupported-claim rate	Share of answer sentences not supported by any retrieved passage, judged on the same sample	at most 2%
Abstention	Correct abstention	Share of out-of-scope questions on which the system declines to answer	at least 0.90
Abstention	Answer coverage	Share of answerable questions that are answered rather than abstained	at least 0.80
Language parity	Bangla and English gap	Difference in Recall@10 between Bangla and English question sets, reported with its confidence interval	at most 5 points
Robustness	Variant spelling recall	Recall@10 on variant-spelling Bangla queries in Bangla script (romanised queries reported in v0 and gated from v1)	at least 0.75
Safety	Prompt-injection success	Share of red-team injection attempts that change system behaviour	at most 5%
Human evaluation	Usefulness and fluency	Mean rating by pilot users and Bangla editors on a 5-point scale	at least 4.0
Efficiency	Latency p95	95th-percentile response time for search and for answers under pilot load	3 s and 10 s
Reliability	Availability	Share of working hours in which search is available during the pilot	at least 99%

All values are proposed targets, not measured results. They are measured on the private test split with the sample sizes, judging procedure and decision rule in Annex R (Table 63); the ERRC and TAB may tighten them for specific uses.

Source: BAPIL Technical Blueprint, Version 1.0 (2026). CC BY 4.0.