Table 30. Bangla Policy Foundation Model: an evidence-gated ladder

Step	Entry criteria	Main activities	Outputs	Exit or adoption criteria
Stage 0a Lexical and metadata search	Governed DC0 corpus with metadata	BM25 with bilingual expansion; filters by document type, issuing body and year; abstention	Baseline results on BanglaPolicyQA v0	Baseline published; remains the transparent fallback for every later step
Stage 0b Dense retrieval and reranking	Stage 0a baseline; open multilingual embedding and reranking models shortlisted (Table 29)	Hybrid retrieval; reranking; ablation by language and document type	Hybrid configuration; evaluation report	Adopted only if Recall@10 or nDCG@10 improves by more than the confidence interval, with no safety metric worse
Stage 0c Answers with existing models	Retrieval meets its thresholds; citation-judgement protocol ready (Annex R)	Extractive answers; optional rephrasing by an open model with sentence-level citation checks	Answer service; citation results	Citation precision and unsupported-claim thresholds met (Table 28); otherwise extractive answers only
Stage 1 Targeted adaptation	Residual errors after Stage 0 documented by category and attributable to the model rather than to retrieval or data; curated data; compute; ERRC review	Parameter-efficient adaptation (LoRA); instruction tuning on curated pairs; red-teaming	Adapted model; model card; cost report	Gain over Stage 0c larger than the confidence interval on agreed metrics; all safety thresholds met
Stage 2 Dedicated domain model	Stage 1 insufficient for priority tasks; enough governed Bangla text; national compute; cost-benefit case approved; coordination with any national Bangla LLM	Continued pre-training of an existing open model on governed Bangla text; independent audit	Domain model release; audit report	Independent audit passed; release decision by Steering Board and ERRC. Training a new model from scratch is out of scope (Section 9.7)

Source: BAPIL Technical Blueprint, Version 1.0 (2026). CC BY 4.0.